


RQ1 Multi-model benchmark on the 98-item expert gold set

Agreement with the gold set



Per-class F_1



Categoriser v1. Intervals are 4,000-replicate bootstrap percentiles. Keyword baseline aligns to 97 of the 98 items (one gold treatment is a parse artifact whose species_id does not match that pipeline's output).